


(a) Same forecast: 8 TP1 GPUs; SLO = 1 s / 20 ms; safety factor = 0.85



(b) Three displayed candidates; clock labels in MHz


	G0	G1	G2	G3	G4	G5	G6	G7	Power W proxy / TPOT ms
A: current M5 + off3	M 2520	M 2520	M 2520	M 2520	M 2520	off	off	off	769.14 214.19 / 16.03
B: reject M4 + off4	M 2520	M 2520	M 2520	M 2520	off	off	off	off	709.55 218.07 / 17.50
C: feasible P2 + D1 + M2 + off3	P 2100	P 2100	D 2520	M 2520	M 2520	off	off	off	708.28 319.17 / 15.54

Screen: proxy ≤ 850 ms; TPOT ≤ 17 ms. C uses the target search without a fixed M floor.

(c) Same-window GPU energy

$E_A = 46.148$ kJ
 $E_C = 42.497 + 0.034 + E_{\text{extra}}$ kJ
 $E_{\text{extra}} \geq 0$: unknown added handoff energy

Normalized cost: A = 100; zoom = 90–100



(d) Hysteresis: compare C with A

Margin $\epsilon = 8\%$
Required gain: 3.692 kJ
Best-case C gain: ≤ 3.617 kJ

$3.617 < 3.692$ kJ
 $92.16 > 92$

C vs A: keep A

Other candidates are not decided.
Synthetic model; no hardware claim.