

The shown safe candidates / TTFT rankings are assumed. Logical slots and bar widths are schematic, not measured or simulated.



Client outputs per request (30 tokens total): green outline = first token; blue outline = subsequent decode token



R6: missing PD first-gap coverage alone does not force M; here M1 is assumed capacity-safe, SLO-safe, incumbent-safe, and best by TTFT.

M2/M3 background load makes them less preferred in this example. R4 waits because earlier P work occupies the available scheduling budget.

K is fully generated in this example. All times are logical slots; no measured latency, throughput, or profile qualification is claimed.